

Hardware-In-the-Loop Robust Co-Design of an Underactuated System

Raphael Stöckner, Shivesh Kumar

Division of Dynamics, Department of Mechanical Engineering, Chalmers University of Technology



CHALMERS
UNIVERSITY OF TECHNOLOGY

Problem

- Co-Design concurrently optimizes system morphology and control, e.g. to give stronger stability guarantees
- Problem: done almost entirely **in simulation**
 - ⚡ Simulation might not capture the system's complexity
 - ⚡ Changing morphology means that different system dynamics may dominate after optimization
 - ⚡ Badly identified dynamics can corrupt the results
- In practice, guarantees might not transfer to the physical system

Closing the Sim-to-Real Gap

- How can **real-world data** be used **in the loop** of the Co-Design process?
- How to minimize the **number of experiments** needed?
- How to **interpret** the sim-to-real **gap** algorithmically?

Parametrizing Robustness

We consider a non-linear torque-controlled inverted pendulum with state $\mathbf{x} = [\theta, \dot{\theta}]^T \in \mathcal{X} \subseteq \mathbb{R}^2$, input $u \in \mathcal{U} \subseteq \mathbb{R}$ and equilibrium $\tilde{\mathbf{x}} = [\pi, 0]^T$. Linearizing the system and applying an LQR controller with gains $\mathbf{Q}, \mathbf{R}$ yields a Lyapunov function $V(\mathbf{x}) = \mathbf{x}^T \mathbf{P} \mathbf{x}$ that guarantees stability of the linearized, unconstrained system.

For the full non-linear closed-loop system, stability only holds within a *region of attraction* (RoA) around $\tilde{\mathbf{x}}$. Since the RoA is generally intractable to compute exactly, it is in practice often approximated as a sublevel set of V :

$$\mathcal{R}_a := \{\mathbf{x} \in \mathcal{X} | \mathbf{x} \rightarrow \tilde{\mathbf{x}} \text{ for } t \rightarrow \infty\} \approx \{\mathbf{x} | V(\mathbf{x}) < \rho\}. \quad (1)$$

This estimate is valid for the largest ρ such that $V(\mathbf{x}) > 0$ and $\dot{V}(\mathbf{x}) < 0$ for all $\mathbf{x} \in \mathcal{R}_a \setminus \{0\}$.

As the RoA gives a formal stability guarantee, its volume can serve as a proxy for robustness [1]: a larger volume means a larger safe region of operation.

Co-Design Formulation

Parameters:

- Morphology: Mass m of the system
- Control: LQR-controller gains $\mathbf{Q}$ and $\mathbf{R}$

Parametrization to avoid over-parametrization and ensure well-definedness of $\mathbf{Q}$ and $\mathbf{R}$:

$$\mathbf{R} = \text{const.} \quad (2)$$

$$\mathbf{Q} = \begin{pmatrix} q_1 & 0 \\ 0 & q_2 \end{pmatrix} = \begin{pmatrix} q_1^{\max} e^{-\bar{q}_1} & 0 \\ 0 & q_2^{\max} e^{-\bar{q}_2} \end{pmatrix}$$

Objective: Maximize the *volume* of the approximated RoA as a proxy for stability and robustness: With $V(\mathbf{x}) = (\mathbf{x} - \tilde{\mathbf{x}})^T \mathbf{P} (\mathbf{x} - \tilde{\mathbf{x}})$ and eigendecomposition $\frac{1}{\rho} \mathbf{P} = \mathbf{W}^T \mathbf{\Lambda} \mathbf{W}$, calculate the volume as

$$\text{vol}(\mathcal{R}_a) = |\det(\sqrt{\mathbf{\Lambda} \mathbf{W}})|^{-1} \frac{\pi}{2}. \quad (3)$$

Optimization Problem: The full optimization problem becomes

$$\max_{m, q_1, q_2} \text{vol}(\mathcal{R}_a) \quad (4a)$$

$$\text{subject to } m \in [m_{\min}, m_{\max}] \quad (4b)$$

$$\bar{q}_i \in \left[0, -\log\left(\frac{q_i^{\min}}{q_i^{\max}}\right)\right], i = 1, 2 \quad (4c)$$


with $\text{vol}(\mathcal{R}_a)$ evaluated with respect to the system dynamics at parameters $m, \bar{q}_1, \bar{q}_2$.

Optimization Algorithm:

- With a *gradient-free, sampling-based* optimization algorithm, the Co-Design formulation could run at *any fidelity level* (i.e. real system, different system models).
- We run it instead at *multiple fidelities concurrently*.
- This allows to combine cheap simulations and more accurate experiments to balance cost versus accuracy

To this end, we adopt a Multi-Fidelity Gaussian Process Regression (MF-GPR) as described in [2]. The idea is to find a surrogate model for the map $f_i : (m, q_1, q_2) \mapsto \text{vol}(\mathcal{R}_a)$ at each fidelity i . The surrogate consists of multiple latent Gaussian Processes and a structure that describes the relationship between them. The step from value-estimation to optimization happens through a so-called acquisition function that selects the next test points that maximizes the information about the maximum of $f_i(m, q_1, q_2)$.

Framework Overview



Results

- Validation: torque-controlled simple pendulum ($m \in [70 \text{ g}, 150 \text{ g}]$, $l = 0.05 \text{ m}$, $u \in \pm 0.04 \text{ N m}$) using the CloudPendulum platform [3]
- 2 fidelities: "real" and "simulation"; simulation uses intentionally inaccurate model (no friction, damping, measurement noise)
- Initialized with 25 simulations and 3 experiments (cost 1 vs. 5) and runs for 52 iterations until the cost budget of 150 is exhausted (199 sims, 28 experiments / ca. 52 min of experiment data)
- Result ($m^* = 70 \text{ g}$, $q_1^* = 7.3 \times 10^{-4}$, $q_2^* = 1.3 \times 10^{-6}$) aligns with physical intuition
- Validation against entirely simulation-based method (using CMA-ES [1]): $\text{vol}(\mathcal{R}_a^{\text{MF-GPR}}) = 47.43$ (ours) vs $\text{vol}(\mathcal{R}_a^{\text{CMA-ES}}) = 28.22$ (theirs)




Figure 1: Estimated Volume of the RoA at "experiment" fidelity for q_1, q_2 with $m = 70 \text{ g}$.




Figure 2: Resulting $\mathcal{R}_a$ as measured on the hardware, using $m = 70 \text{ g}$.

- [1] Lasse Jennings Maywald et al. "Co-optimization of Acrobot Design and Controller for Increased Certifiable Stability". In: *2022 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. 2022, pp. 2636–2641.
- [2] Shion Takeno et al. "Multi-fidelity Bayesian optimization with max-value entropy search and its parallelization". In: *International Conference on Machine Learning*. PMLR. 2020, pp. 9334–9345.
- [3] Shivesh Kumar. "Swinging pendulums on the cloud: Digitalization of simulation & experimental infrastructure for feedback-based active learning". In: *Chalmers Konferens om undervisning och lärande (KUL2025)*. 2025.



Learn more! Paper, CloudPendulum, etc.